



Introduction

The Credit Card Clients Dataset is a comprehensive financial dataset that provides crucial insights into credit card usage patterns and client behavior. Featuring a diverse array of variables, including demographic information, credit limits, payment history, and bill amounts, this dataset facilitates in-depth analysis of credit card performance and risk assessment. With a focus on client credit behavior and repayment tendencies, it serves as a valuable resource for financial institutions, researchers, and data scientists seeking to enhance their understanding of consumer credit dynamics. The dataset's richness enables the development of predictive models to assess creditworthiness and optimize financial strategies for risk management.

1 import Necessary Library

```
In [69]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

2 import Dataset

```
In [70]: df = pd.read_csv("/kaggle/input/default-of-credit-card-clients-dataset/UCI_Credit_Card.csv")
```

3 Data Analysis

```
In [71]: df.head()
```

```
Out[71]: ID LIMIT BAL SEX EDUCATION MARRIAGE AGE PAY 0 PAY 2 PAY 3 PAY 4 BILL AMT4
```

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	...	BILL_AMT
0	1	20000.0	2	2	1	24	2	2	-1	-1	...	0.0
1	2	120000.0	2	2	2	26	-1	2	0	0	...	3272.0
2	3	90000.0	2	2	2	34	0	0	0	0	...	14331.0
3	4	50000.0	2	2	1	37	0	0	0	0	...	28314.0
4	5	50000.0	1	2	1	57	-1	0	-1	0	...	20940.0

5 rows × 25 columns




In [72]:

```
df.tail()
```

Out[72]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	...	BILL_AMT
29995	29996	220000.0	1	3	1	39	0	0	0	0	...	
29996	29997	150000.0	1	3	2	43	-1	-1	-1	-1	...	
29997	29998	30000.0	1	2	2	37	4	3	2	-1	...	
29998	29999	80000.0	1	3	1	41	1	-1	0	0	...	
29999	30000	50000.0	1	2	1	46	0	0	0	0	...	

5 rows × 25 columns



In [73]:

```
df.columns
```

Out[73]:

```
Index(['ID', 'LIMIT_BAL', 'SEX', 'EDUCATION', 'MARRIAGE', 'AGE', 'PAY_0',  
      'PAY_2', 'PAY_3', 'PAY_4', 'PAY_5', 'PAY_6', 'BILL_AMT1', 'BILL_AMT2',  
      'BILL_AMT3', 'BILL_AMT4', 'BILL_AMT5', 'BILL_AMT6', 'PAY_AMT1',  
      'PAY_AMT2', 'PAY_AMT3', 'PAY_AMT4', 'PAY_AMT5', 'PAY_AMT6',  
      'default.payment.next.month'],  
      dtype='object')
```

In [74]:

```
df.shape
```

Out[74]: (30000, 25)

In [75]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 30000 entries, 0 to 29999  
Data columns (total 25 columns):  
#   Column              Non-Null Count  Dtype  
---  ---  
0   ID                   30000 non-null  int64  
1   LIMIT_BAL            30000 non-null  float64  
2   SEX                  30000 non-null  int64  
3   EDUCATION            30000 non-null  int64  
4   MARRIAGE              30000 non-null  int64  
5   AGE                   30000 non-null  int64  
6   PAY_0                 30000 non-null  int64  
7   PAY_2                 30000 non-null  int64  
8   PAY_3                 30000 non-null  int64  
9   PAY_4                 30000 non-null  int64
```

```
10 PAY_5 30000 non-null int64
11 PAY_6 30000 non-null int64
12 BILL_AMT1 30000 non-null float64
13 BILL_AMT2 30000 non-null float64
14 BILL_AMT3 30000 non-null float64
15 BILL_AMT4 30000 non-null float64
16 BILL_AMT5 30000 non-null float64
17 BILL_AMT6 30000 non-null float64
18 PAY_AMT1 30000 non-null float64
19 PAY_AMT2 30000 non-null float64
20 PAY_AMT3 30000 non-null float64
21 PAY_AMT4 30000 non-null float64
22 PAY_AMT5 30000 non-null float64
23 PAY_AMT6 30000 non-null float64
24 default.payment.next.month 30000 non-null int64
dtypes: float64(13), int64(12)
memory usage: 5.7 MB
```

In [76]:

df.describe()

Out[76]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000
mean	15000.500000	167484.322667	1.603733	1.853133	1.551867	35.485500	-0
std	8660.398374	129747.661567	0.489129	0.790349	0.521970	9.217904	1
min	1.000000	10000.000000	1.000000	0.000000	0.000000	21.000000	-2
25%	7500.750000	50000.000000	1.000000	1.000000	1.000000	28.000000	-1
50%	15000.500000	140000.000000	2.000000	2.000000	2.000000	34.000000	0
75%	22500.250000	240000.000000	2.000000	2.000000	2.000000	41.000000	0
max	30000.000000	1000000.000000	2.000000	6.000000	3.000000	79.000000	8

8 rows × 25 columns

In [77]:


df.corr()

Out[77]:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	
	ID	1.000000	0.026179	0.018497	0.039177	-0.029079	0.018678 -0
	LIMIT_BAL	0.026179	1.000000	0.024755	-0.219161	-0.108139	0.144713 -0
	SEX	0.018497	0.024755	1.000000	0.014232	-0.031389	-0.090874 -0
	EDUCATION	0.039177	-0.219161	0.014232	1.000000	-0.143464	0.175061 0
	MARRIAGE	-0.029079	-0.108139	-0.031389	-0.143464	1.000000	-0.414170 0
	AGE	0.018678	0.144713	-0.090874	0.175061	-0.414170	1.000000 -0
	PAY_0	-0.030575	-0.271214	-0.057643	0.105364	0.019917	-0.039447 1
	PAY_2	-0.011215	-0.296382	-0.070771	0.121566	0.024199	-0.050148 0
	PAY_3	-0.018494	-0.286123	-0.066096	0.114025	0.032688	-0.053048 0
	PAY_4	-0.002735	-0.267460	-0.060173	0.108793	0.033122	-0.049722 0
	PAY_5	-0.022199	-0.249411	-0.055064	0.097520	0.035629	-0.053826 0

PAY_6	-0.020270	-0.235195	-0.044008	0.082316	0.034345	-0.048773	C
BILL_AMT1	0.019389	0.285430	-0.033642	0.023581	-0.023472	0.056239	C
BILL_AMT2	0.017982	0.278314	-0.031183	0.018749	-0.021602	0.054283	C
BILL_AMT3	0.024354	0.283236	-0.024563	0.013002	-0.024909	0.053710	C
BILL_AMT4	0.040351	0.293988	-0.021880	-0.000451	-0.023344	0.051353	C
BILL_AMT5	0.016705	0.295562	-0.017005	-0.007567	-0.025393	0.049345	C
BILL_AMT6	0.016730	0.290389	-0.016733	-0.009099	-0.021207	0.047613	C
PAY_AMT1	0.009742	0.195236	-0.000242	-0.037456	-0.005979	0.026147	-C
PAY_AMT2	0.008406	0.178408	-0.001391	-0.030038	-0.008093	0.021785	-C
PAY_AMT3	0.039151	0.210167	-0.008597	-0.039943	-0.003541	0.029247	-C
PAY_AMT4	0.007793	0.203242	-0.002229	-0.038218	-0.012659	0.021379	-C
PAY_AMT5	0.000652	0.217202	-0.001667	-0.040358	-0.001205	0.022850	-C
PAY_AMT6	0.003000	0.219595	-0.002766	-0.037200	-0.006641	0.019478	-C
default.payment.next.month	-0.013952	-0.153520	-0.039961	0.028006	-0.024339	0.013890	C

25 rows × 25 columns



In [78]:

df.ndim

Out[78]: 2

4 Data cleaning and Preprocessing

In [79]:

df.isnull().sum()

```
Out[79]: ID                0
LIMIT_BAL                0
SEX                      0
EDUCATION                0
MARRIAGE                 0
AGE                      0
PAY_0                    0
PAY_2                    0
PAY_3                    0
PAY_4                    0
PAY_5                    0
PAY_6                    0
BILL_AMT1                0
BILL_AMT2                0
BILL_AMT3                0
BILL_AMT4                0
BILL_AMT5                0
BILL_AMT6                0
PAY_AMT1                 0
PAY_AMT2                 0
PAY_AMT3                 0
PAY_AMT4                 0
PAY_AMT5                 0
PAY_AMT6                 0
```

```
default.payment.next.month    0  
dtype: int64
```

```
In [80]: df.isna().empty
```

```
Out[80]: False
```

```
In [81]: df.isna().sum()
```

```
Out[81]: ID                0  
LIMIT_BAL                0  
SEX                      0  
EDUCATION                0  
MARRIAGE                 0  
AGE                      0  
PAY_0                    0  
PAY_2                    0  
PAY_3                    0  
PAY_4                    0  
PAY_5                    0  
PAY_6                    0  
BILL_AMT1                0  
BILL_AMT2                0  
BILL_AMT3                0  
BILL_AMT4                0  
BILL_AMT5                0  
BILL_AMT6                0  
PAY_AMT1                 0  
PAY_AMT2                 0  
PAY_AMT3                 0  
PAY_AMT4                 0  
PAY_AMT5                 0  
PAY_AMT6                 0  
default.payment.next.month    0  
dtype: int64
```

```
In [82]: df['default.payment.next.month'].value_counts()
```

```
Out[82]: default.payment.next.month  
0      23364  
1       6636  
Name: count, dtype: int64
```

5| Data visualisation

EDA (Exploratory Data Analysis)

5.1 displot

```
In [83]: sns.distplot(df, kde = False, bins=30)  
plt.show()
```


/tmp/ipykernel_42/305216106.py:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

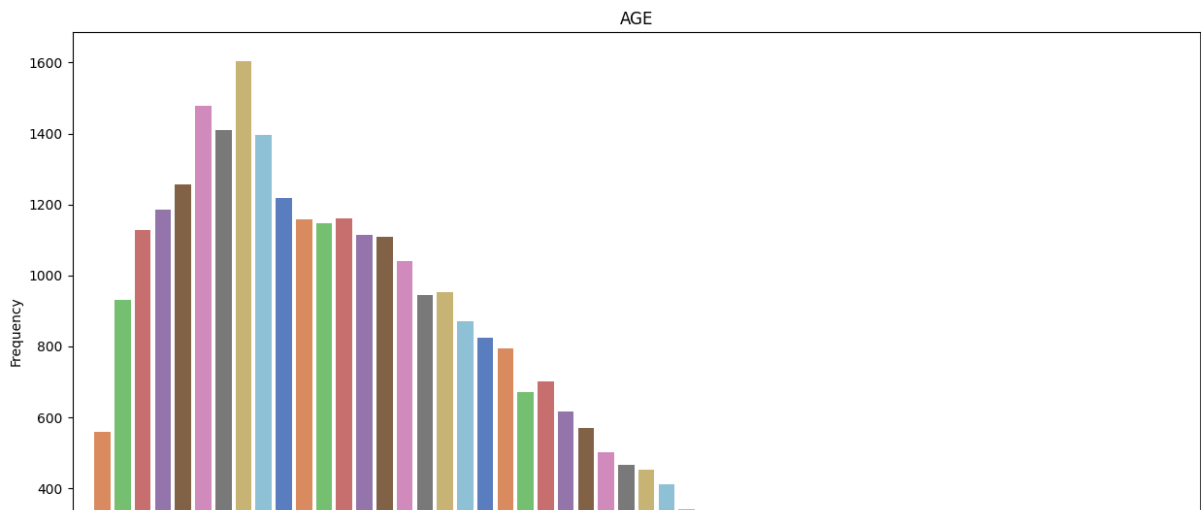
```
sns.distplot(df, kde = False, bins=30)
```




5.2 countplot

In [84]:


```
# sepal_length
plt.figure(figsize=(15, 8))
sns.countplot(x='AGE', data=df, palette='muted')
plt.title('AGE')
plt.xlabel('Lenght')
plt.ylabel('Frequency')
plt.xticks(rotation=45, ha='right')
plt.show()
```



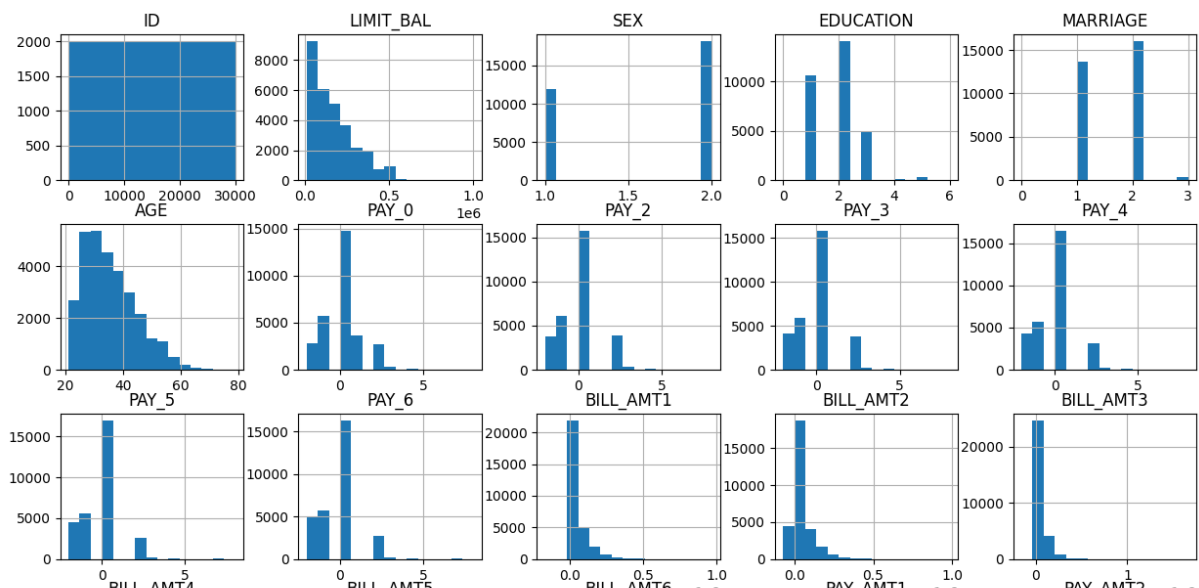



```
In [85]: sns.countplot(data=df, x='AGE', y=None, hue=None,
                      order=None, hue_order=None, orient=None,
                      color=None, palette=None, saturation=.75,
                      width=.8, dodge=True, ax=None)
```

Out[85]: <Axes: xlabel='AGE', ylabel='count'>



```
In [86]: df.hist(figsize=(15,12),bins = 15)
plt.title("Features Distribution")
plt.show()
```





6 | Split the Dataset

```
In [87]: from sklearn.model_selection import train_test_split
```

```
In [88]: X = df.drop(['default.payment.next.month'], axis = 1)
```

```
In [89]: y = df['default.payment.next.month']
```

```
In [90]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

```
In [91]: X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[91]: ((22500, 24), (7500, 24), (22500,), (7500,))
```

7 | PCA (Principal Component Analysis)

```
In [92]: from sklearn.decomposition import PCA
```

```
In [93]: pca = PCA(n_components=2)
```

```
In [94]: X_pca = pca.fit_transform(X)
```

```
In [95]: X_pca[0]
```

```
Out[95]: array([-166511.13375728, -75548.89621183])
```

```
In [96]: print("Explained Variance Ratio:")
print(pca.explained_variance_ratio_)
```

```
Explained Variance Ratio:
[0.60943217 0.2948671 ]
```

```
In [97]: from sklearn.preprocessing import MinMaxScaler
```



```
scaler = MinMaxScaler()
X_ft = scaler.fit_transform(X)
```

In [98]: `X_ft[0]`

```
Out[98]: array([0.00000000e+00, 1.01010101e-02, 1.00000000e+00, 3.33333333e-01,
 3.33333333e-01, 5.17241379e-02, 4.00000000e-01, 4.00000000e-01,
 1.00000000e-01, 1.00000000e-01, 0.00000000e+00, 0.00000000e+00,
 1.49981727e-01, 6.91643226e-02, 8.67228923e-02, 1.60137756e-01,
 8.06480880e-02, 2.60978723e-01, 0.00000000e+00, 4.09081976e-04,
 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00])
```

In [99]: `X_ft[1]`

```
Out[99]: array([3.33344445e-05, 1.11111111e-01, 1.00000000e+00, 3.33333333e-01,
 6.66666667e-01, 8.62068966e-02, 1.00000000e-01, 4.00000000e-01,
 2.00000000e-01, 2.00000000e-01, 2.00000000e-01, 4.00000000e-01,
 1.48892434e-01, 6.78575089e-02, 8.78171337e-02, 1.63219937e-01,
 8.40739510e-02, 2.63484742e-01, 0.00000000e+00, 5.93732912e-04,
 1.11602161e-03, 1.61030596e-03, 0.00000000e+00, 3.78310691e-03])
```



Machine Learning Algorithm

Algorithm

(1) KNN

In [100... `from sklearn.neighbors import KNeighborsClassifier`
`from sklearn.metrics import accuracy_score`
`from sklearn.model_selection import train_test_split`

In [101... `knn_classifier = KNeighborsClassifier(n_neighbors=3)`

In [102... `knn_classifier.fit(X_train, y_train)`

Out[102... `KNeighborsClassifier(n_neighbors=3)`

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [103... `train_predictions = knn_classifier.predict(X_train)`
`train_accuracy1 = accuracy_score(y_train, train_predictions)`

In [104... `test_predictions = knn_classifier.predict(X_test)`
`test_accuracy1 = accuracy_score(y_test, test_predictions)`

In [105... `print(f"Training Accuracy: {train_accuracy1}")`

```
print(f"Training Accuracy: {train_accuracy1} ")
print(f"Testing Accuracy: {test_accuracy1}")
```

Training Accuracy: 0.8444

Testing Accuracy: 0.7334666666666667

(2) Naive Bayes classifier

```
In [106... from sklearn.naive_bayes import GaussianNB
from sklearn.naive_bayes import BernoulliNB
from sklearn.naive_bayes import MultinomialNB

from sklearn import metrics
```

```
In [107... # GaussianNB
```

```
In [108... G_classifier = GaussianNB()
```

```
In [109... G_classifier.fit(X_train, y_train)
```

Out[109... GaussianNB()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [110... train_predictions = G_classifier.predict(X_train)

train_accuracy21 = accuracy_score(y_train, train_predictions)
```

```
In [111... test_predictions = G_classifier.predict(X_test)

test_accuracy21 = accuracy_score(y_test, test_predictions)
```

```
In [112... print(f"Training Accuracy: {train_accuracy21}")
print(f"Testing Accuracy: {test_accuracy21}")
```

Training Accuracy: 0.37684444444444444

Testing Accuracy: 0.3817333333333333

```
In [113... # BernoulliNB
```

```
In [114... B_classifier = BernoulliNB()
```

```
In [115... B_classifier.fit(X_train, y_train)
```

Out[115... BernoulliNB()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

nbviewer.org.

```
In [116... train_predictions = B_classifier.predict(X_train)

train_accuracy22 = accuracy_score(y_train, train_predictions)
```

```
In [117... test_predictions = G_classifier.predict(X_test)

test_accuracy22 = accuracy_score(y_test, test_predictions)
```

```
In [118... print(f"Training Accuracy: {train_accuracy22}")
print(f"Testing Accuracy: {test_accuracy22}")
```

Training Accuracy: 0.7717777777777778

Testing Accuracy: 0.3817333333333333

(3) Decision Tree

```
In [119... from sklearn.tree import DecisionTreeClassifier
```

```
In [120... clf = DecisionTreeClassifier()
```

```
In [121... clf.fit(X_train, y_train)
```

Out[121... DecisionTreeClassifier()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [122... train_predictions = clf.predict(X_train)

train_accuracy3 = accuracy_score(y_train, train_predictions)
```

```
In [123... test_predictions = clf.predict(X_test)

test_accuracy3 = accuracy_score(y_test, test_predictions)
```

```
In [124... print(f"Training Accuracy: {train_accuracy3}")
print(f"Testing Accuracy: {test_accuracy3}")
```

Training Accuracy: 1.0

Testing Accuracy: 0.7298666666666667

(4) Random Forest

```
In [125... from sklearn.ensemble import RandomForestClassifier
```

```
In [126... rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
In [127... rf_classifier.fit(X_train, y_train)
```

```
Out[127... RandomForestClassifier(random_state=42)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [128... train_predictions = rf_classifier.predict(X_train)

train_accuracy4 = accuracy_score(y_train, train_predictions)
```

```
In [129... test_predictions = rf_classifier.predict(X_test)

test_accuracy4 = accuracy_score(y_test, test_predictions)
```

```
In [130... print(f"Training Accuracy: {train_accuracy4}")
print(f"Testing Accuracy: {test_accuracy4}")
```

Training Accuracy: 1.0
Testing Accuracy: 0.8172

(5) Boosting Algorithm

```
In [131... from sklearn.ensemble import AdaBoostClassifier
```

```
In [132... base_classifier = DecisionTreeClassifier(max_depth=1)
```

```
In [133... adaboost_classifier = AdaBoostClassifier(base_classifier, n_estimators=50, random_state=42)
```

```
In [134... adaboost_classifier.fit(X_train, y_train)
```

```
Out[134... AdaBoostClassifier(estimator=DecisionTreeClassifier(max_depth=1),
                    random_state=42)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [135... train_predictions = adaboost_classifier.predict(X_train)

train_accuracy5 = accuracy_score(y_train, train_predictions)
```

```
In [136... test_predictions = adaboost_classifier.predict(X_test)

test_accuracy5 = accuracy_score(y_test, test_predictions)
```

In [137]...

```
print(f"Training Accuracy: {train_accuracy5}")
print(f"Testing Accuracy: {test_accuracy5}")
```

Training Accuracy: 0.8185777777777777

Testing Accuracy: 0.8181333333333334

(6). Logistic Regression

In [138]...

```
from sklearn import linear_model
```

In [139]...

```
lrg = linear_model.LogisticRegression()
```

In [140]...

```
lrg.fit(X_train, y_train)
```

/opt/conda/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

Out[140]...

```
n_iter_i = _check_optimize_result(
LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [141]...

```
train_predictions = lrg.predict(X_train)
train_accuracy7 = accuracy_score(y_train, train_predictions)
```

In [142]...

```
test_predictions = lrg.predict(X_test)
test_accuracy7 = accuracy_score(y_test, test_predictions)
```

In [143]...

```
print(f"Training Accuracy: {train_accuracy7}")
print(f"Testing Accuracy: {test_accuracy7}")
```

Training Accuracy: 0.7772888888888889

Testing Accuracy: 0.7830666666666667

(7).Linear Regression

In [144]...

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
```

In [145]...

```
model = LinearRegression()
```

In [146]...

```
model.fit(X_train, y_train)
```

```
model.fit(X_train, y_train)
```

Out[146... LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [147...

```
train_predictions = clf.predict(X_train)
```

```
train_accuracy8 = accuracy_score(y_train, train_predictions)
```

In [148...

```
test_predictions = clf.predict(X_test)
```

```
test_accuracy8 = accuracy_score(y_test, test_predictions)
```

In [149...

```
print(f"Training Accuracy: {train_accuracy8}")  
print(f"Testing Accuracy: {test_accuracy8}")
```

Training Accuracy: 1.0

Testing Accuracy: 0.7298666666666667

(8).Gradient Boosting Machines (GBM)

In [150...

```
from sklearn.ensemble import GradientBoostingClassifier
```

In [151...

```
model = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)
```

In [152...

```
model.fit(X_train, y_train)
```

Out[152... GradientBoostingClassifier(random_state=42)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [153...

```
train_predictions = model.predict(X_train)
```

```
train_accuracy9 = accuracy_score(y_train, train_predictions)
```

In [154...

```
test_predictions = model.predict(X_test)
```

```
test_accuracy9 = accuracy_score(y_test, test_predictions)
```

In [155...

```
print(f"Training Accuracy: {train_accuracy9}")  
print(f"Testing Accuracy: {test_accuracy9}")
```

Training Accuracy: 0.8265333333333333

Testing Accuracy: 0.8216

(6).SVM

```
In [156... from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
```

```
In [157... scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [158... svm_classifier = SVC(kernel='linear', C=1.0)
```

```
In [159... svm_classifier.fit(X_train, y_train)
```

```
Out[159... SVC(kernel='linear')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [160... train_predictions = svm_classifier.predict(X_train)

train_accuracy6 = accuracy_score(y_train, train_predictions)
```

```
In [161... test_predictions = svm_classifier.predict(X_test)

test_accuracy6 = accuracy_score(y_test, test_predictions)
```

```
In [162... print(f"Training Accuracy: {train_accuracy6}")
print(f"Testing Accuracy: {test_accuracy6}")
```

Training Accuracy: 0.8096888888888889

Testing Accuracy: 0.8089333333333333

Decision Tree, Random Forest, SVM, Gradient Boosting Machines (GBM), Algorithm is the best accuracy

accuracy = 0.8




10 | Hierarchical Clustering

```
In [163... from scipy.cluster.hierarchy import linkage, dendrogram, fcluster
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
```

```
In [164... scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [165... linkage_matrix = linkage(X_scaled, method='ward')
```

```
In [166... plt.figure(figsize=(12, 6))
dendrogram(linkage_matrix, labels=df['default.payment.next.month'].values, orientation='t')
plt.title('Hierarchical Clustering Dendrogram')
plt.xlabel('default.payment.next.month')
plt.ylabel('Distance')
plt.show()
```



THANK YOU!

